

# From Idea to Hardware

High-Level Synthesis with Haskell

- Yifan Yu
- Mark Arrumm
- Rupesh Majhi

# Introduction suomeksi

- Hankkeessamme olemme tutkineet erilaisia tapoja suunnitella digitaalista järjestelmää.
- Vertaamme Verilog-, C-, ja Haskell-ohjelmointikieliä, mutta Haskell-kieli on pääasiallinen tutkimuskohde.
- Esitämme Haskell-kielisen toteutuksen ja teemme havaintoesityksen.

# Table of Contents

- Chapter 1: History of Functional Programming and Haskell
- Chapter 2: 3 Ways of Describing Hardware
- Chapter 3: Our Idea -> Our Hardware
- Chapter 4: Live demonstration

Computers compute. But what is computation?



# Abacus vs CPU

From a low-level point of view, an abacus does the same thing as a CPU. So hypothetically and theoretically, can we run our Windows software on top of an abacus?



# Let's talk about history!



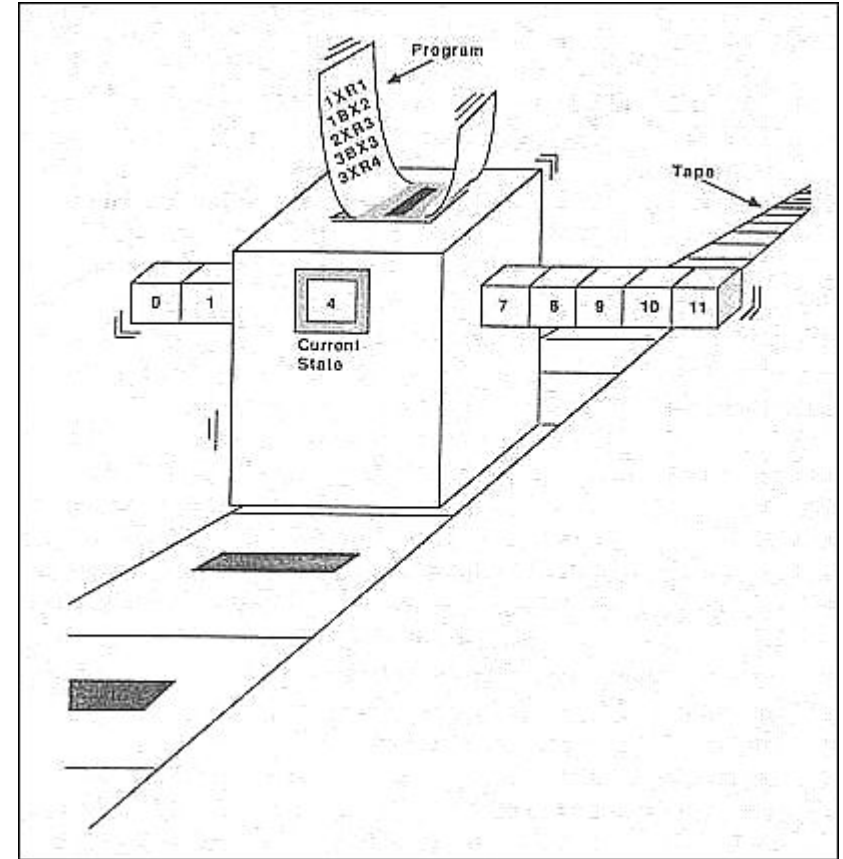
# Computability

- **David Hilbert (1900)**: proof that the arithmetic is *consistent* – free of any internal contradictions (2nd problem out of 23)
- **Kurt Gödel (1931)**: it is not! (Gödel's incompleteness theorems)
  - A function is computable *if and only if* it is a *general recursive function*
- **Alan Turing (1936)**: a function is computable *if and only if* it is computable by a Turing Machine
- **Alonzo Church (1936)**: a function is computable *if and only if* it can be written down as a term of the  $\lambda$ -calculus

# Turing machine

A Turing machine is an abstract machine that manipulates symbols on a strip of tape according to a table of rules.

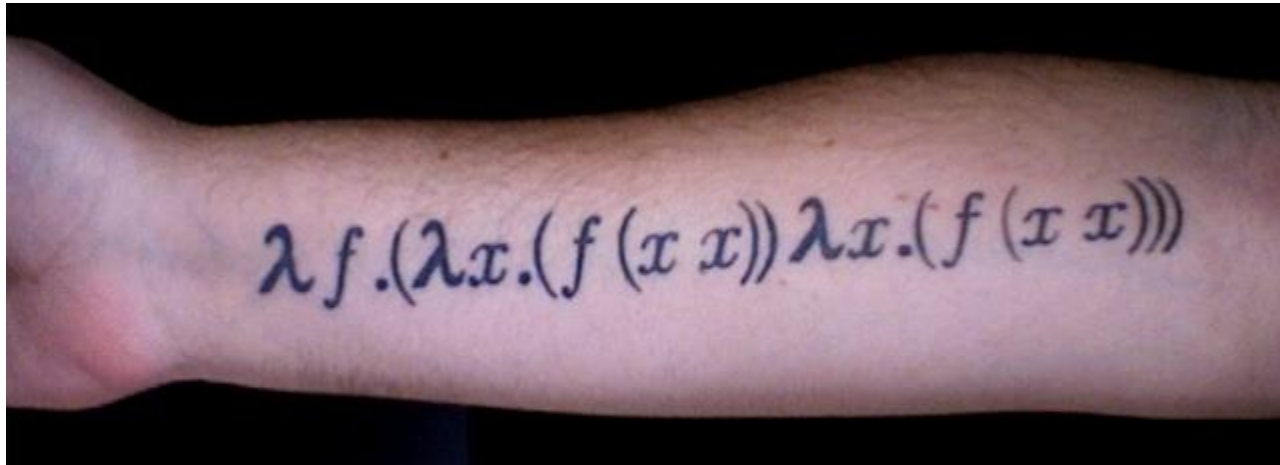
A system is **Turing complete** if it can solve any problem that a Turing machine can.





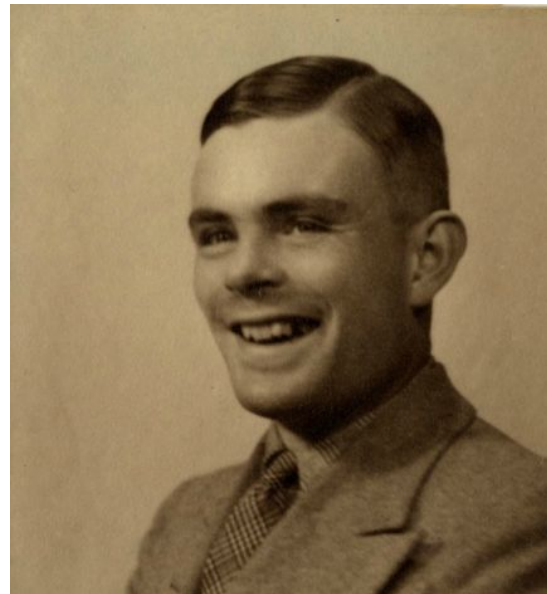
# $\lambda$ -calculus

- A  $\lambda$ -term is either:
  - A variable:  $x$
  - Lambda abstraction:  $\lambda x. t$
  - Application:  $t x$



# Church-Turing thesis

- They are equivalent.
- *In other words:* if there is an algorithm to solve a problem, then it can be solve with a Turing Machine or the  $\lambda$ -calculus
- *In other words:* a Turing Machine or the  $\lambda$ -calculus can simulate any mechanical computer
- *In other words:* everything feasibly computable in the physical world is feasibly computable by a Turing Machine or the  $\lambda$ -calculus



# FP is born

- Based on  $\lambda$ -calculus as theoretical framework
  - Lisp in MIT, late 1950s -> lots of dialects (e.g., Scheme, Clojure)  
The final lines of NASA's LISP based launch script :  
))))))))))))))))))))))  
))))))))))))))))))))))))))
  - APL, early 1960s
  - ML, 1970s -> Standard ML, OCaml, F#
  - Haskell, Erlang, 1980s
  - Scala (multi-paradigm), 2004
- Mainstream languages are adding FP features: Java, C++ added lambda functions, etc.

# Haskell features

- Immutable variables
- Pure functions (no side effects)
- Higher-order functions
- Lazy evaluation
- Pattern matching
- Advanced type system
  - Full type inference
  - Strong & static typing
  - Type classes
  - Type polymorphism



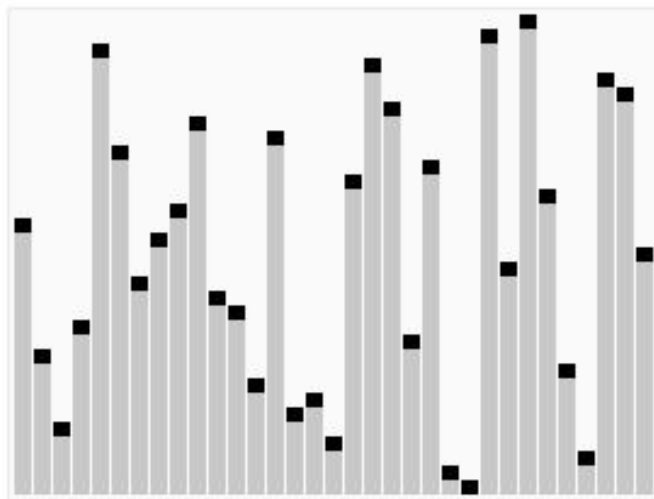
Haskell Curry

# A Taste of Haskell

## Quicksort

```
quickSort [] = []
```

```
quickSort (x:xs) = quickSort (filter (<x) xs) ++ [x] ++ quickSort (filter (>=x) xs)
```



# In comparison with C

```
void quicksort(int x[10],int first,int last)
{
    int pivot,j,temp,i;

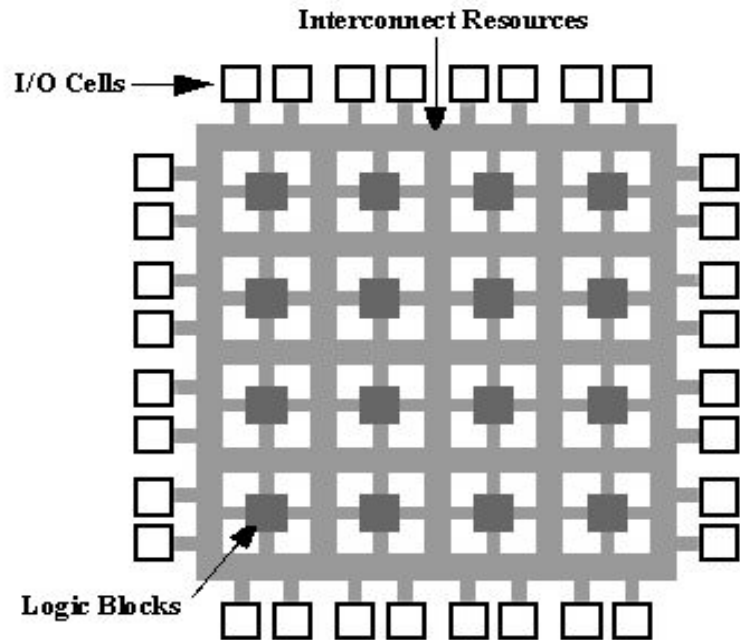
    if (first < last) {
        pivot = first;
        i     = first;
        j     = last;

        while (i < j) {
            while (x[i] <= x[pivot] && i < last)
                i++;
            while (x[j] > x[pivot])
                j--;
```

```
        if (i < j) {
            temp = x[i];
            x[i] = x[j];
            x[j] = temp;
        }
    }

    temp      = x[pivot];
    x[pivot] = x[j];
    x[j]      = temp;
    quicksort(x, first, j-1);
    quicksort(x, j+1 , last);
}
}
```

# FPGA - Field Programmable Gate Array

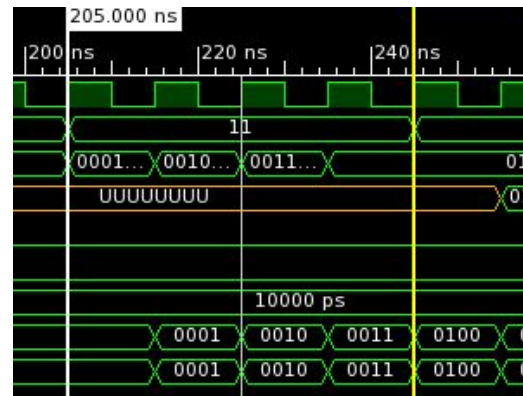
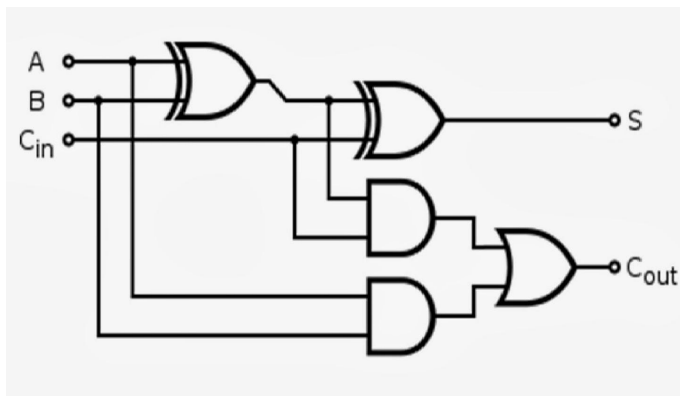




# 3 Ways of Describing Hardware

- What most people are using: *Verilog, VHDL*
- What some people have attempted: *HLS with C/C++/SystemC*
- What we are trying: *Clash (Haskell based, designed by researchers in University of Twente, Netherlands)*

```
always@(posedge clk)
  if (clear == 1)
    r_out <= 0;
  else
    r_out <= r_out + a * b;
assign out = r_out;
```



# Example 1: a simple OR gate

```
module or_test(  
    input [1:0] sw,  
    output led  
);
```

```
    assign led =  
        sw[0] | sw[1];
```

```
endmodule
```

*Verilog*

```
void or(int *a, int *b,  
        int *out)  
{  
    *out = *a | *b;  
}
```

*C*

```
or a b = a .|. B
```

*Haskell*

## Example 1: a simple OR gate

```
or a b = a .|. B
```

Not only is this simpler, it also has a new feature:

```
or :: Bits a => a -> a -> a
```

Type polymorphism makes it *generic*, so that it can be applied to different types.

(not possible with Verilog or C)

## Example 2: a MAC unit (multiply-accumulate)

```
module mac(  
    input clk,  
    input clear,  
    input [15:0] a,  
    input [15:0] b,  
    output [15:0] out)  
    reg [15:0] state;  
  
    always@(posedge clk)  
        if (clear == 1)  
            state <= 0;  
        else  
            state <= state + a * b;  
    assign out = state;  
endmodule
```

*Verilog*

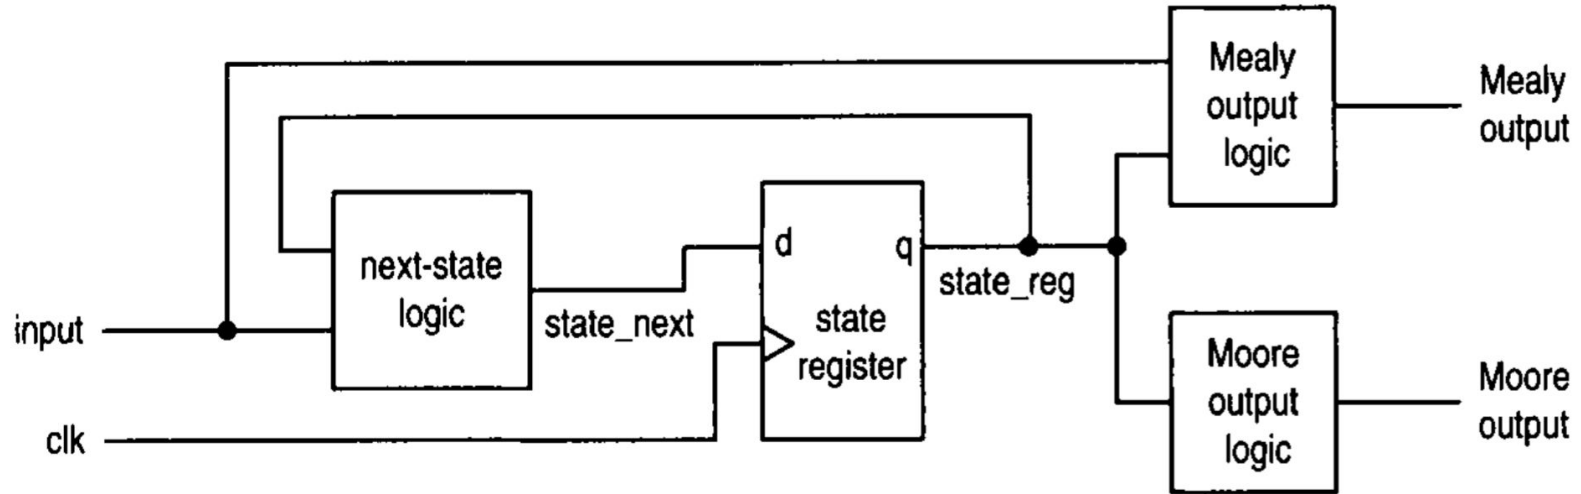
```
void macc(int a, int b,  
          int *accum, bool clr)  
{  
    static acc_reg = 0;  
    if (clr)  
        acc_reg = 0;  
    acc_reg += a * b;  
    *accum = acc_reg;  
}
```

*C*

```
macT acc input = (f acc input, acc)  
    where  
        f a (x, y) = a + x * y  
  
mac = mealy macT (0, 0)
```

*Haskell*

# Mealy machine vs Moore machine



# The higher-order **mealy** function

```
mealy :: (s -> i -> (s, o))      state -> in -> (newstate, out)
      -> s                        initial state
      -> (Signal i -> Signal o) synchronous sequential function
```

The MAC unit we saw earlier:

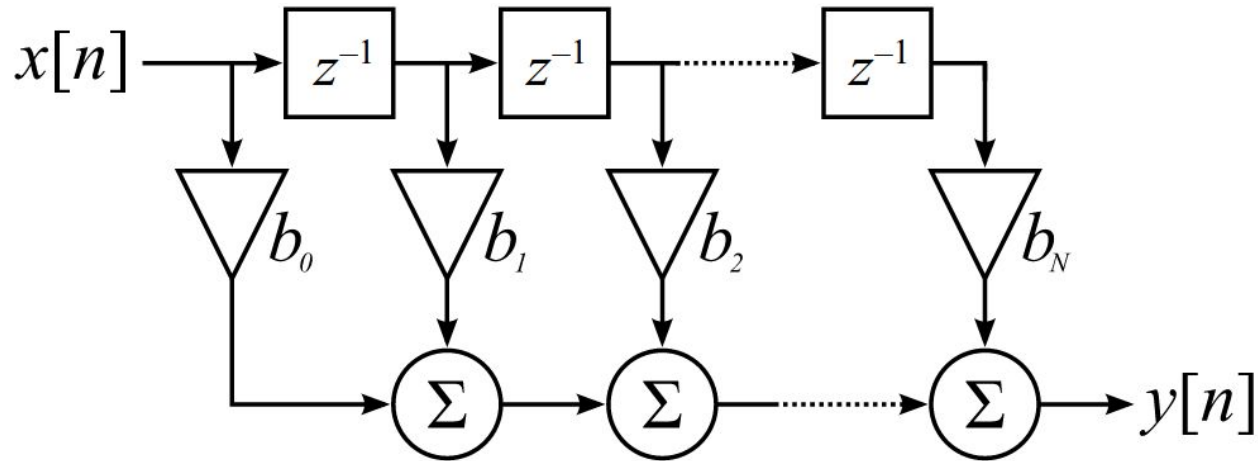
```
macT acc input = (f acc input, acc)
  where
    f a (x, y) = a + x * y

mac = mealy macT (0, 0)
```

Higher-order function **abstracts away** some of the common patterns, so you can construct any Mealy machine with a pure state transition function.

## Example 3: an FIR filter (finite impulse response)

- Verilog & C implementation are too long to show here
- FIR filter definition: the dot-product of a set of filter coefficients and a window over the input



## Example 3: an FIR filter (finite impulse response)

- Verilog & C implementation are too long to show here
- FIR filter definition: the dot-product of a set of filter coefficients and a window over the input

### *Haskell*

```
fir coeffs input = dotp coeffs (window input)
  where
    dotp as bs = sum (zipWith (*) as bs)
```

Notice how closely the code matches the definition!



# A hierarchy of abstractions

**Software design:** machine code -> assembly code -> C -> OOP, FP, etc

**Hardware design:** HDL -> HLS ( C -> FP)

```
01110011 01    push    ebp
01100101 01    mov     ebp, esp
01101000 01    movzx   ecx, [ebp]
01100100 01    pop     ebp
01100100 01    movzx   dx, cl
01110010 01    lea     eax, [edx]
01110100 01    add     eax, edx
01100001 01    shl     eax, 2
01101001 01    add     eax, edx
01101101 01    shr     eax, 8
01101101 01    sub     cl, al
00100000 01    shr     cl, 1
01110011 01    add     al, cl
```

```
-----Monte Carlo
while(t1<=tmax)
{
    fprintf(fp, "%f\t%d\n", t1, l);
    t1+=dt;
    for(j=0; j<=n0; j++)
    {
        r=rand()%1000;
        r=r/1000;
        if(r<=0.001) n0--;
        if(n0<0) goto l1;
    }
}
-----
fclose(fp);
```

```
-- Type annotation (optional)
fib :: Int -> Integer

-- With self-referencing data
fib n = fibs !! n
      where fibs = 0 : scanl (
                -- 0,1,1,2,3,5,...
            ) 1 fibs

-- Same, coded directly
fib n = fibs !! n
      where fibs = 0 : 1 : next
                next (a : t@(b:_))
```

# Circuit vs Haskell

## *Circuit*

- Combinational circuits
- Undefined behavior OK if you don't use it
- D flip-flop “delays” a signal by one cycle
- Self-reference

## *Haskell*

- Pure functions
- Undefined value ok if you don't use it (lazy evaluation)
- (:) "delays" everything in a list by one element
- Self-reference

# Now, let's design something in Haskell!

- It has to be sufficiently sophisticated
- At the same time, simple enough for us to pick up CλaSH and finish the design in just few weeks
- It has to reflect issues we need to address in real-world projects
- It has to be fun

**A Brainfuck CPU**

# What the Brainfuck?

- A minimalism programming language with an ungodly name
- Has only 8 commands: > < + - . , [ ]
- Everything else gets ignored
- Operates on an array of bytes
- **It is fully Turing-complete**
- In face, it is very close to a Turing Machine

Show me the code!

```
+++++++ [ >++++++>++++++>+++>+  
<<<<-] >++.>+.+++++. .+++.>++.<<++++  
+++++++ .>.+++ .----- .-----.>+  
.>.
```

That's the “Hello, World!” Program

# So... it's pretty useless, why make a CPU for it?

- Software interpreter is trivial to implement, CPU is harder
- Simple instructions, but demonstrates the main principles of CPU design
- We need to address common issues
  - Architecture choice
  - Instruction timing
  - Caches
- Perfect for testing:
  - the capabilities of CλaSH
  - the feasibility of describing hardware with Haskell
- It sounds like fun!

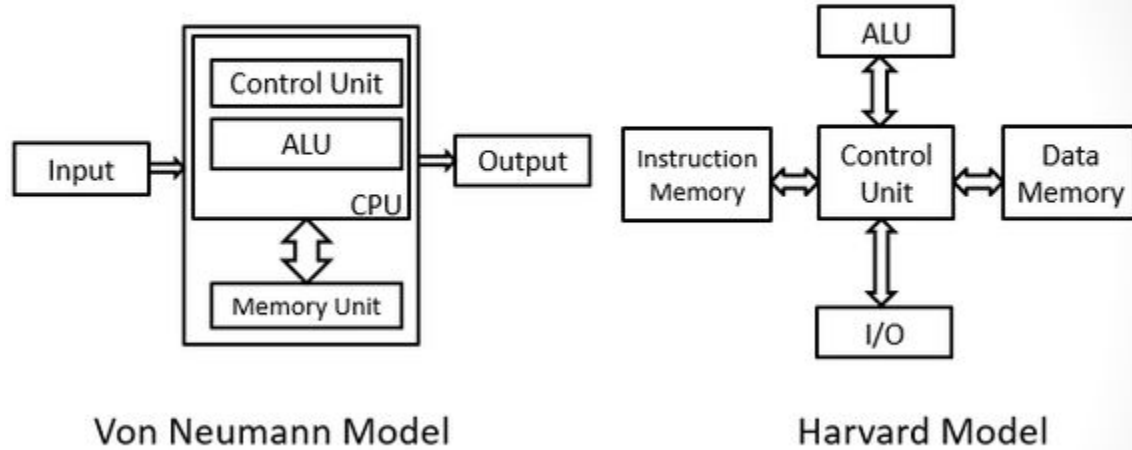
# What's new in our design?

- The idea of a Brainfuck CPU is not new. Different approaches:
  - VHDL
  - Lava (Haskell-based DSL)
  - Python + MyHDL
  - Directly with logic gates
- Our CλaSH-based design compiles directly from Haskell to Verilog
  - Enables testing and debugging with REPL
  - High-level constructs are used such as **mealy**
- Directly programmable through UART
- Counts clock cycles and displays on seven-segment display
- Has a debugging mode which enables clock-by-clock debugging
- A number of optimizations are exploited

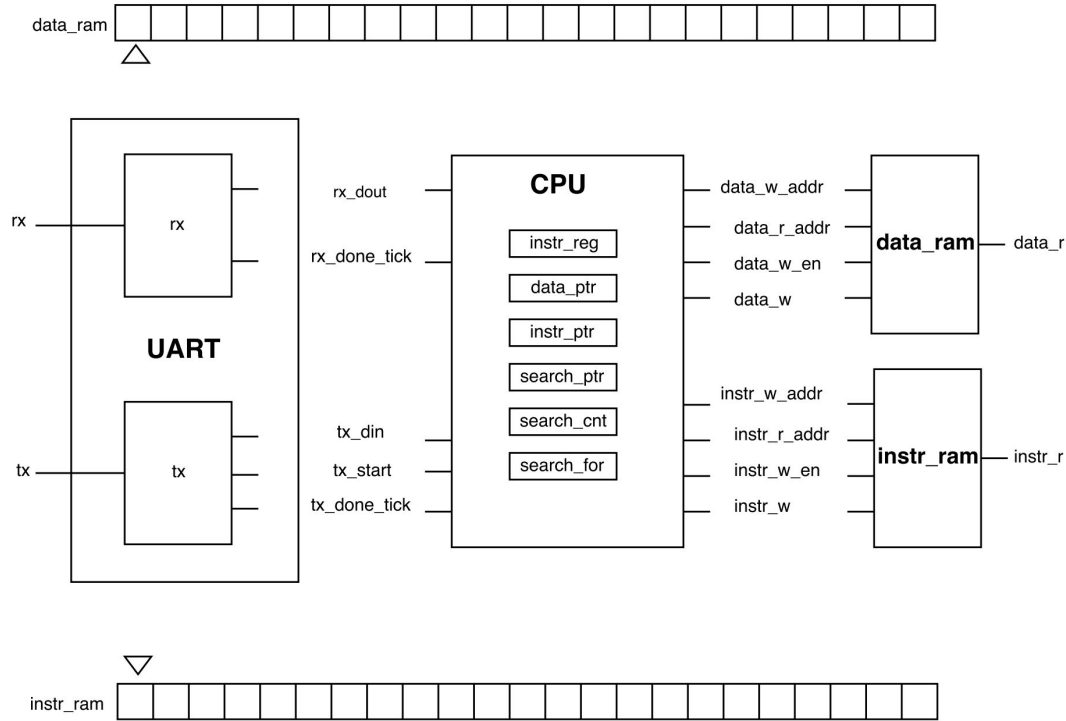


# Design: architecture

## Von Neumann vs. Harvard architectures



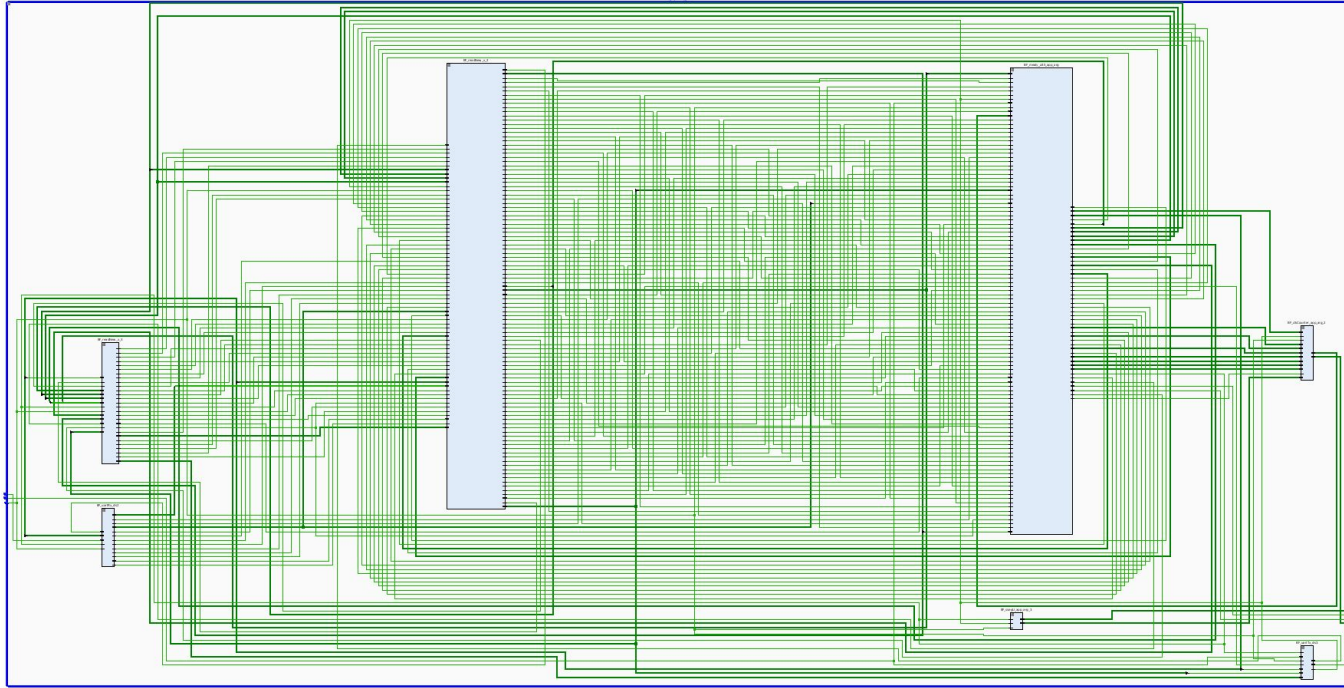
# Design: Version 1



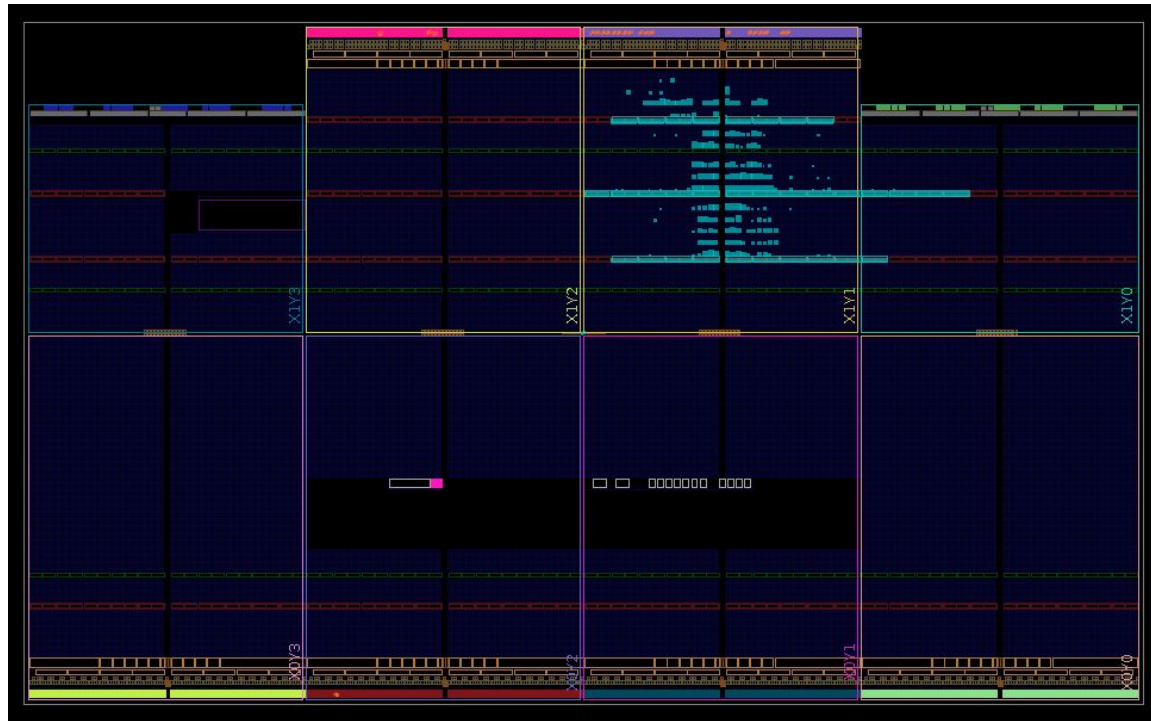
# Results

- Version 1: **648** lines of Haskell (including SSeg, UART) generated **3036** lines of Verilog
- Version 1 is able to run test programs, *nicely and slowly*
- More optimizations can be exploited
  - Jump address pre-indexing
  - Common instruction patterns can be reduced to a single instruction
  - Pipelining
- More updates on GitHub: <https://github.com/aufheben/clash-bfcpu>

# Synthesized Schematic



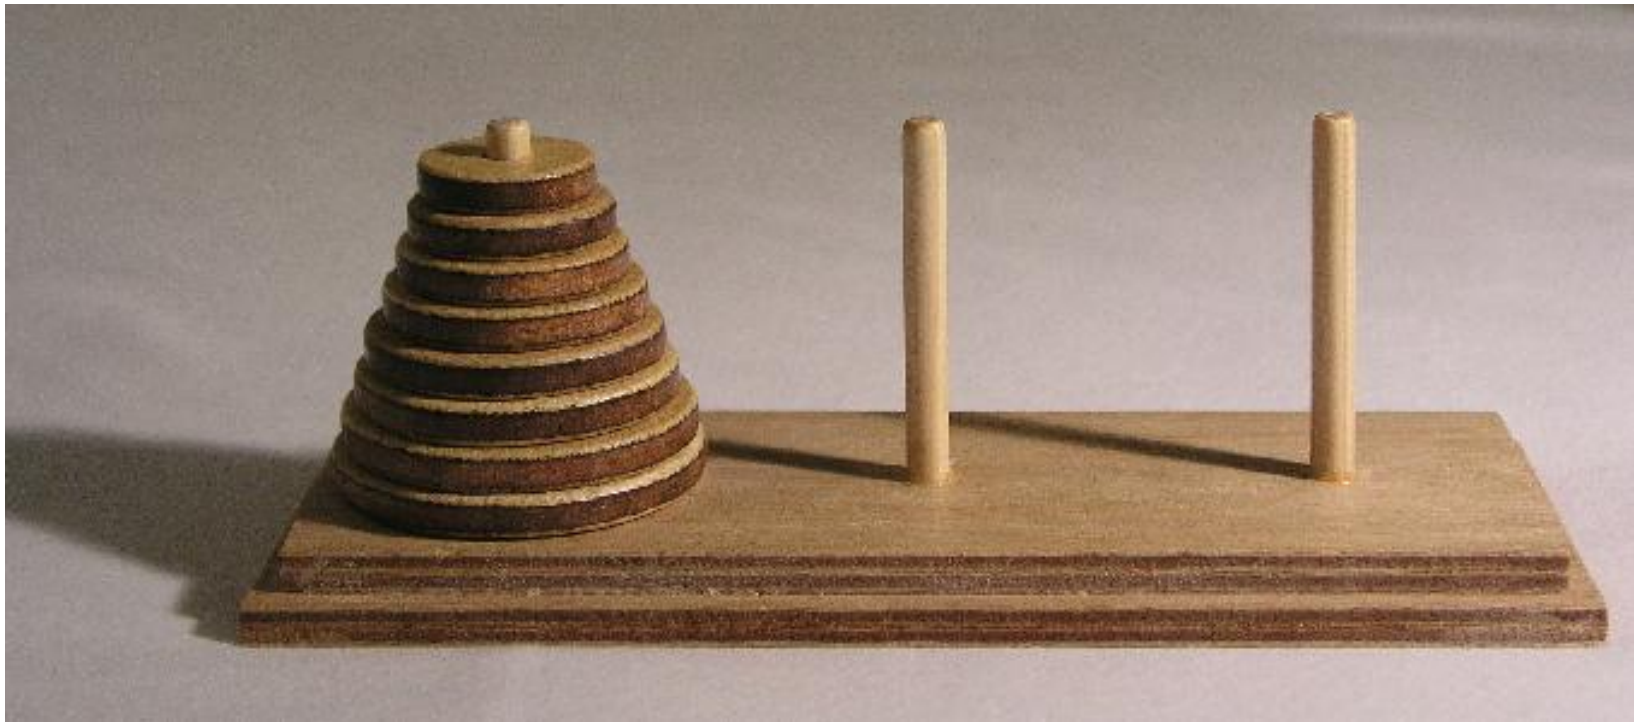
# Implementation



Now, live demonstration!

+++++++ [ >+++++++>+++++++>+++>+  
<<<<- ] >+ . >+ . +++++ . . + . >+ . <<++++  
+++++++ . > . + . - - - - - . - - - - - . >+  
. > .

# Tower of Hanoi (click to download)



# [Screenshot From Live Demo]

Towers of Hanoi in Brainf\*ck  
Written by Clifford Wolf <<http://www.clifford.at/bfcpu/>>

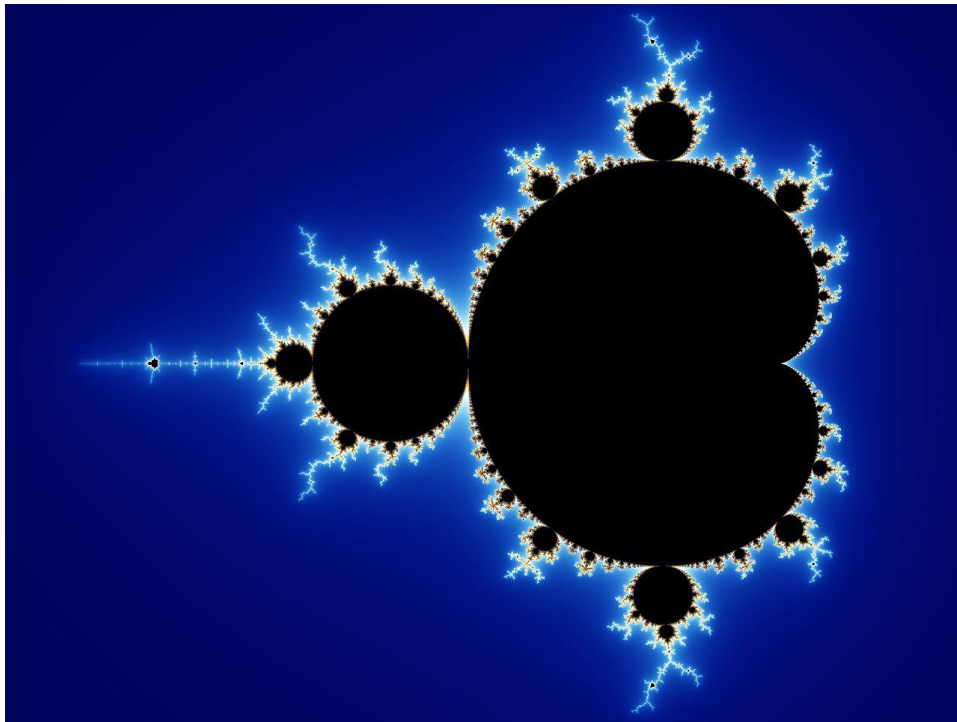
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx  
xxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxxx

xxxxxxxxxxxxxxxxxxxxxxxxx  
xxxxxxxxxxxxxxxxxxxxxxxxx  
xxxxxxxxxxxxxxxxxxxxxxxxx

  xXx  
  xxxxxx  
  xxxxxxxxx  
  xxxxxxxxxx



# Mandelbrot set (click to download)



[Screenshot From Live Demo]

[illegible]

Try: [Brainfuck Interpreter Written in Brainfuck](#)



An optimized & pipelined **CPU** of the  
*Brainfuck*  
programming language, designed with  
**Clash** and **Haskell**  
from scratch, implemented on  
**FPGA**

Wed 14.12  
11:00 - 11:30  
B407

Yifan Yu  
Mark Arrumm  
Rupesh Majhi

[github.com/aufheben/clash-bfcpu](https://github.com/aufheben/clash-bfcpu)

*"It's just stunning!"* - Kai Lindgren  
*"I'm astonished."* - Timo Kasurinen

